

```

+-----+
|      CSE 521      |
| PROJECT 1: THREADS |
| DESIGN DOCUMENT   |
+-----+

```

---- GROUP ----

>> Fill in the names and email addresses of your group members.

```

Ved Harish Valsangkar    <vedharis@buffalo.edu>
Gursimran Singh         <gursimr2@buffalo.edu>
Harshal Ganesh Jagtap    <harshalg@buffalo.edu>

```

---- PRELIMINARIES ----

>> If you have any preliminary comments on your submission, notes for the
>> TAs, or extra credit, please give them here.

>> Please cite any offline or online sources you consulted while
>> preparing your submission, other than the Pintos documentation, course
>> text, lecture notes, and course staff.

```

https://www.youtube.com/watch?v=npQF28g6s_k
https://www.youtube.com/watch?v=S12qx1DwjVk

```

```

ALARM CLOCK
=====

```

---- DATA STRUCTURES ----

>> A1: Copy here the declaration of each new or changed `struct' or
>> `struct' member, global or static variable, `typedef', or
>> enumeration. Identify the purpose of each in 25 words or less.

Adding a new variable to the thread struct to hold the timestamp at which
time the thread must be taken out of blocked_list.

```

struct thread
{
    ...
    int64_t sleep_wt;           /* Wake time stamp for looping through
blocked list. */
    ...
};

```

A new list created to hold the blocked threads thread.c

```

static struct list blocked_list;

```

Lock for ready_list created in thread.h

```

struct lock rl_lock;

```

Lock for blocked_list created in thread.h

```

struct lock bl_lock;

```

Comparator for comparing the wake-up time of the threads for sorting
blocked_list.

```
bool timer_priority_comparator(const struct list_elem *a,
                               const struct list_elem *b,
                               void *aux)
```

---- ALGORITHMS ----

>> A2: Briefly describe what happens in a call to timer_sleep(),
>> including the effects of the timer interrupt handler.

By our current algorithm, We first calculate the wait time for our current thread, which will set the variable sleep_wt to the expected tick stamp when the process should be woken up. Then we acquire lock for our blocked queue and push the current process in the blocked queue. The blocked_list will always be sorted by wake time. Then we release the lock for blocked queue, and call the thread_block() method, which changes the current status of our thread from THREAD_RUNNING to THREAD_BLOCKED and call the schedule() method which reschedules our thread.

>> A3: What steps are taken to minimize the amount of time spent in
>> the timer interrupt handler?

Since, the blocked_list is in sorted state we need check the whether the head element needs to be unblocked. If so, we unblock all process whose timer has elapsed. If not, no time is spent on further processing.

---- SYNCHRONIZATION ----

>> A4: How are race conditions avoided when multiple threads call
>> timer_sleep() simultaneously?

As per our current mechanism, the process of putting a thread into the blocked_list requires the blocked_list to acquire lock and will release the lock after successful addition to the blocked_list before thread_block() calls schedule(). Hence, there won't be any new process in the running state before thread is put into blocked_list. That means that no two threads will try to access blocked_list simultaneously. Hence, race conditions are avoided.

>> A5: How are race conditions avoided when a timer interrupt occurs
>> during a call to timer_sleep()?

In timer_sleep(), we have used a lock, bl_lock, to lock editing of the blocked_list. This lock prevents anyone trying to access the blocked_list, even timer_interrupt

---- RATIONALE ----

>> A6: Why did you choose this design? In what ways is it superior to
>> another design you considered?

To get rid of the busy waiting loop, we need to put the running thread into blocked_list instead of ready_list. By blocking the running thread instead of putting it into ready_list, we are ensuring that the thread doesn't get scheduled before its waiting time completion by putting it into the blocked_list.

Using this design, we could easily understand the flow of the threads and visualize them as if on a timeline. This implementation is simple and concise, as compared to other possibilities like locking the threads using individual locks.

PRIORITY SCHEDULING
=====

---- DATA STRUCTURES ----

>> B1: Copy here the declaration of each new or changed 'struct' or 'struct' member, global or static variable, 'typedef', or enumeration. Identify the purpose of each in 25 words or less.

Adding a new variable, `donated_priority`, to the thread struct to hold the pointer to the priority variable which can be re-assigned to transfer priority. Another variable, `locks_held`, has been added to track how many threads have requested lock held by the current thread.

```
struct thread
{
    ...
    int * donated_priority;
    int locks_held;
    ...
};
```

Comparator for comparing the priority of the threads for sorting `ready_list`.

```
bool priority_comparator(const struct list_elem *a,
                        const struct list_elem *b,
                        void *aux)
```

>> B2: Explain the data structure used to track priority donation.
>> Use ASCII art to diagram a nested donation. (Alternately, submit a .png file.)

---- ALGORITHMS ----

>> B3: How do you ensure that the highest priority thread waiting for a lock, semaphore, or condition variable wakes up first?

The `ready_list` is already sorted in a descending order according to priority. Thus higher priority threads are closer to the head. Two threads with equal priority are treated on an FCFS basis. Thus `ready_list` priority requirement is satisfied. The `blocked_list` is sorted based on the time stamp at which the thread is expected to be woken up in an ascending order. Hence, all the necessary time stamps can be woken up if required at any given tick. The sorting also ensures no thread is woken up before their wait time is elapsed regardless of the priority.

>> B4: Describe the sequence of events when a call to `lock_acquire()` causes a priority donation. How is nested donation handled?

When `lock_acquire` is called on already acquired lock, we compare the correct owner's priority with the current process. If less, we assign the current thread's `donated_priority` pointer to the owner's `donated_priority`. We increment the current lock owner's attribute `locks_held` to keep track of number of dependent locks.

>> B5: Describe the sequence of events when `lock_release()` is called on a lock that a higher-priority thread is waiting for.

When `lock_release` is called, we decrement attribute `locks_held`. If `locks_held` becomes 0, we redirect `donated_priority` to point to our current priority and release the lock.

---- SYNCHRONIZATION ----

>> B6: Describe a potential race in `thread_set_priority()` and explain
 >> how your implementation avoids it. Can you use a lock to avoid
 >> this race?

---- RATIONALE ----

>> B7: Why did you choose this design? In what ways is it superior to
 >> another design you considered?

This was a simplistic design imitating `alarm_clock` where where the list
 sorting would be done by a similar comparator. Such a implementation
 can now be expanded to suit any other need.

ADVANCED SCHEDULER =====

---- DATA STRUCTURES ----

>> C1: Copy here the declaration of each new or changed `'struct'` or
 >> `'struct'` member, global or static variable, `'typedef'`, or
 >> enumeration. Identify the purpose of each in 25 words or less.

An array of lists shall be used to store multiple queues, each consisting
 of threads with equal priorities. The total number of such queues
 would be 64 corresponding to 64 priority levels.

```
static struct list mlfqs_table;
```

A global variable, `load_avg`, is used to store the average load calculated
 by the formula given in slides.

```
static int32_t load_avg;
```

The `round_robin_flag` is set everytime a time slice is completed to indicate
 recalculation of priority of the current thread.

```
static bool round_robin_flag;
```

This flag is set just before the timer interrupt calls `thread_yield()` to
 indicate `thread_yield()` has to evaluate per-tick operations.

```
bool intr_flag;
```

---- ALGORITHMS ----

>> C2: Suppose threads A, B, and C have nice values 0, 1, and 2. Each
 >> has a `recent_cpu` value of 0. Fill in the table below showing the
 >> scheduling decision and the priority and `recent_cpu` values for each
 >> thread after each given number of timer ticks:

timer ticks	recent_cpu			priority			thread to run
	A	B	C	A	B	C	
0	0	0	0	63	61	59	A
4	4	0	0	62	61	59	A
8	8	0	0	61	61	59	A
12	12	0	0	60	61	59	B
16	12	4	0	60	60	59	B
20	12	8	0	60	59	59	A

24	16	8	0	59	59	59	A
28	20	8	0	58	59	59	B
32	20	12	0	58	58	59	C
36	20	12	4	58	58	58	C

>> C3: Did any ambiguities in the scheduler specification make values in the table uncertain? If so, what rule did you use to resolve them? Does this match the behavior of your scheduler?

Considerations:

- If two processes including the running process end up having same priority, then the running process is given priority.
- If two processes other than running process end up having same priority, then there is no specific information given to tackle such a clash. We propose using nice values as tie-breakers with lower nice values getting higher priority.

>> C4: How is the way you divided the cost of scheduling between code inside and outside interrupt context likely to affect performance?

To compensate for the difficulties of accessing variables inside and outside interrupt context, we had to use flag and global variables. Owing to this, programmer oversight and byzantine errors could cause untracable, unreproducible runtime errors. Also, for any buildup over this project, certain aspects of this code would need to be changed to meet new set of requirements.

---- RATIONALE ----

>> C5: Briefly critique your design, pointing out advantages and disadvantages in your design choices. If you were to have extra time to work on this part of the project, how might you choose to refine or improve your design?

The idea behind an array of lists is that we have a fixed number of priorities and hence we require a fixed number of queues to hold the threads according to their corresponding priority. Hash maps in theory could have helped us in reducing fetching time for these queues, but we found the implementation more complex than what was required for this task.

>> C6: The assignment explains arithmetic for fixed-point math in detail, but it leaves it open to you to implement it. Why did you decide to implement it the way you did? If you created an abstraction layer for fixed-point math, that is, an abstract data type and/or a set of functions or macros to manipulate fixed-point numbers, why did you do so? If not, why not?

We referred an online video resource to understand an implementation of fixed-point math in practice. We attempted to adapt logic written for cpp to native c which caused errors because of difficulty in understanding of operator precedence. We used a combination of the online resource and the lecture slides to create our own functions. We tried to use pre-compiler commands for simplicity and clarity in coding.

SURVEY QUESTIONS

=====

Answering these questions is optional, but it will help us improve the course in future quarters. Feel free to tell us anything you want--these questions are just to spur your thoughts. You may also choose to respond anonymously in the course evaluations at the end of

the quarter.

>> In your opinion, was this assignment, or any one of the three problems
>> in it, too easy or too hard? Did it take too long or too little time?

It was difficult to adapt to pointer based procedural programming language
especially managing pointer arithmetic and scope.

>> Did you find that working on a particular part of the assignment gave
>> you greater insight into some aspect of OS design?

>> Is there some particular fact or hint we should give students in
>> future quarters to help them solve the problems? Conversely, did you
>> find any of our guidance to be misleading?

A crash course in pointers and assignment would be beneficial.

>> Do you have any suggestions for the TAs to more effectively assist
>> students, either for future quarters or the remaining projects?

>> Any other comments?